

To Ponder

- ❑ Pick 5 cards
- ❑ Prepare to be astonished...

Security: Cryptography

Computer Science and Engineering ■ College of Engineering ■ The Ohio State University

Lecture 26

Some High-Level Goals

- Confidentiality
 - Non-authorized users have limited access
- Integrity
 - Accuracy/correctness/validity of data
- Availability
 - No down-time or disruptions
- Authenticity
 - Agents are who they claim to be
- Non-repudiation
 - A party to a transaction can not later deny their participation

Methods of Attack

- ❑ Target people (“social engineering”)
 - Phishing: email, phone, surveys, ...
 - Baiting: click & install, physical media, ...
- ❑ Target software (“exploits”)
 - Unpatched OS, browser, programs
 - Buffer overflow
 - Code injection and cross-site scripting
- ❑ Target channel (“man-in-the-middle”)
 - Eavesdropping
 - Masquerading, tampering, replay

Cryptography

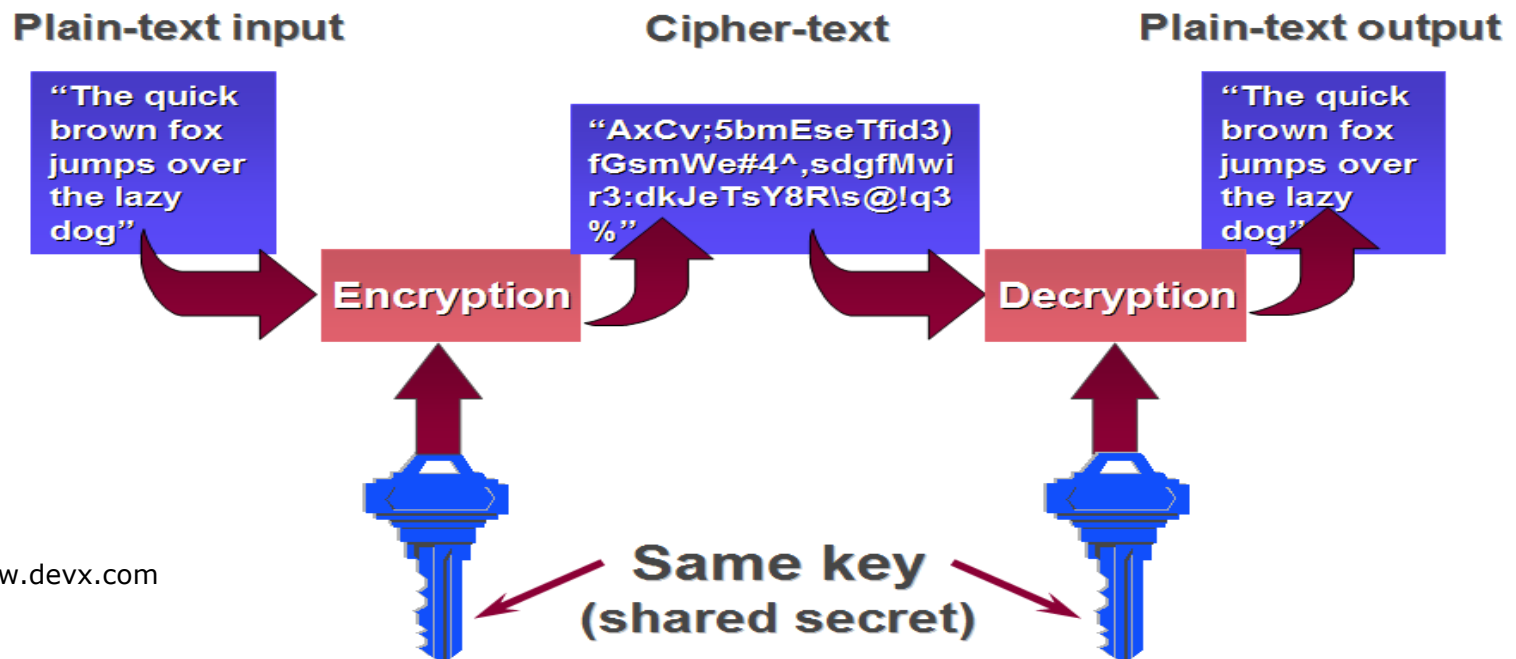
- Etymology (Greek)
 - *kryptos*: hidden or secret
 - *grapho*: write
- Basic problem:
 - 2 agents (traditionally “**A**lice” and “**B**ob”)
 - A & B want to exchange private messages
 - Channel between A & B is not secure (“Eve” is eavesdropping)
- Solution has other applications too
 - Protect *stored* data (e.g. on disk, or in cloud)
 - Digital *signatures* for non-repudiation
 - Secure *passwords* for authentication

Core Idea: The Secret

- Alice & Bob share some *secret*
 - Secret can not be the message itself
 - Secret used to protect arbitrary messages
- Crude analogy: a padlock
 - Copies of the physical key are the secret
 - Alice puts message in box and locks it
 - Bob unlocks box and reads message
- But real channels are bit streams
 - Eve can see the bits!
 - Message must be garbled in some way
 - Secret is strategy for garbling/degarbling


Protecting Messages

- ❑ Alice garbles (encrypts) the message
- ❑ Sends the encrypted cipher-text
- ❑ Bob knows how to degarble (decrypt) cipher-text back into plain-text



Encryption/Decryption Function

Computer Science and Engineering ■ The Ohio State University



$$\begin{aligned} E: P &\rightarrow Q \\ D: Q &\rightarrow P \end{aligned}$$

$$\begin{aligned} E(m) &= c \\ D(c) &= m \\ \text{i.e. } D &= E^{-1} \end{aligned}$$

Note: often $P = Q$
So E is a *permutation*

Families of Encryption Functions

- Each pair of agents needs their own E
 - Many E's (& corresponding D's) needed
- But good E's are hard to invent
- Solution: design one (good) E, which is parameterized by a number
 - That is, have a huge *family* of E's:
 $E_0, E_1, E_2, \dots E_K$
 - Secret: which E_i is used (i is the **key**)

Classic Example: Caesar Cipher

- Shift each letter by x positions in alphabet
 - *E.g.* $x = 3$
 $a \rightarrow d, b \rightarrow e, c \rightarrow f, d \rightarrow g, e \rightarrow h, \dots$
 - The key is x
- Encode a string character-by-character
 - For $m = \text{"hello world"}$, $E_3(m) = \text{"khoor zruog"}$
- Questions:
 - What is P (set of plaintext messages)?
 - What is Q (set of ciphertext messages)?
 - How many different ciphers?
 - Is this a strong or weak cipher?


Generalized Caesar Cipher

- Generalization: arbitrary mapping
 - *E.g.* The qwerty shift
 $a \rightarrow s, b \rightarrow n, c \rightarrow v, d \rightarrow f, e \rightarrow r, \dots$
 - For $m = \text{"hello world"}$,
 $E(m) = \text{"jraap eptaf"}$
 - 26! possible ciphers... that's a lot!
 - Approximately 4×10^{26}
 - There are $\sim 10^{18}$ nanoseconds/century
- Weakness?

Generalized Caesar Cipher

- Generalization: arbitrary mapping
 - *E.g.* The qwerty shift
 $a \rightarrow s, b \rightarrow n, c \rightarrow v, d \rightarrow f, e \rightarrow r, \dots$
 - For $m = \text{"hello world"}$,
 $E(m) = \text{"jraap eptaf"}$
 - 26! possible ciphers... that's a lot!
 - Approximately 4×10^{26}
 - There are $\sim 10^{18}$ nanoseconds/century
- Weakness?
 - In English text, letters appear in predictable ratios
 - From enough ciphertext, can infer E

Frequency Analysis



Leon Battista Alberti

Computer Science and Engineering ■ The Ohio State University



WW II: Enigma Machine

Computer Science and Engineering ■ The Ohio State University



Polyalphabetic Cipher

- Alberti's idea: Use different E_i 's within the same message
 - $E(\text{"hello world"}) = E_a(\text{"h"})E_b(\text{"e"})E_c(\text{"l"})E_d(\text{"l"})E_e(\text{"o"})\dots$
- Alice & Bob need to agree on the *sequence* of E 's to use
- Claude Shannon proved that this method is perfectly secure (1949)
 - Precise information-theoretic meaning
 - Known as a *one-time pad*

One-Time Pad

- Message is a sequence of bits

$m_0 \ m_1 \ m_2 \ m_3 \ m_4 \ m_5 \ m_6 \dots$

- One-time pad is *random* bit sequence

$x_0 \ x_1 \ x_2 \ x_3 \ x_4 \ x_5 \ x_6 \dots$

- E is bit-wise XOR operation, \oplus

- Cipher text is

$m_0 \oplus x_0 \ m_1 \oplus x_1 \ m_2 \oplus x_2 \ m_3 \oplus x_3 \ m_4 \oplus x_4 \ m_5 \oplus x_5 \ m_6 \oplus x_6 \dots$

- Problem: Pad is long and cannot be re-used (hence cumbersome to share)
- “Solution”: pseudo-random sequence, generated from a seed (the key)

Comparison: Stream vs Block

Stream Cipher


- Encrypts bit-by-bit
- $|P| = |Q| = 2$
- Few choices for E (roughly 2)
- Message can have any length

Block Cipher


- Encrypts a fixed-length (k -bit) sequence
- $|P| = |Q| = 2^k$
- Many choices for E (roughly 2^k !)
- Padding added s.t. $|m| \bmod k = 0$

Limitation of Fixed Block Size

- ❑ Message can be longer than block size
- ❑ Reuse same E for each block?
 - Danger: Frequency analysis vulnerability
 - Don't do this (for multiblock messages)!

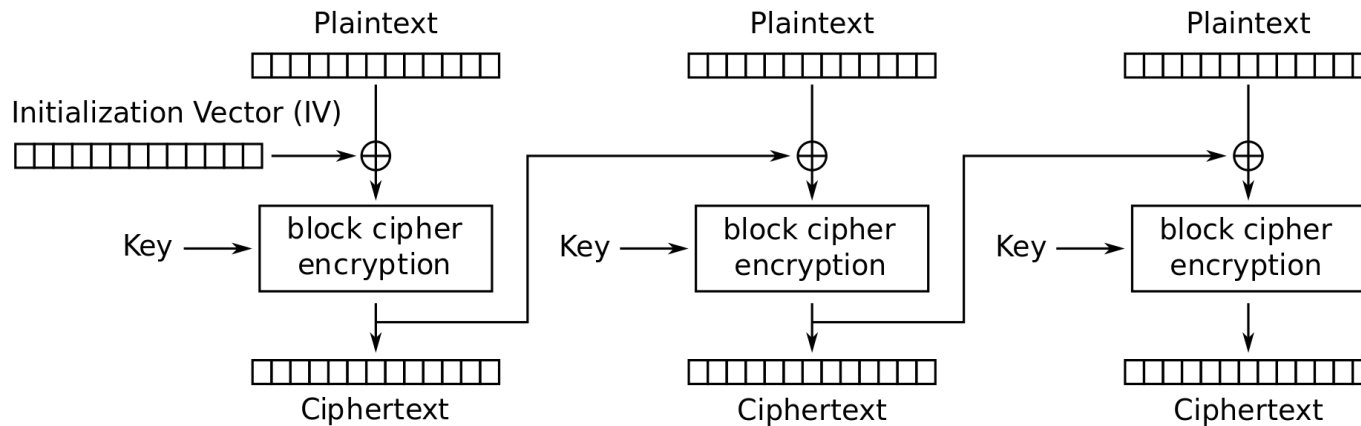


Electronic Codebook (ECB) mode encryption



Solution: Initialization Vector


- ❑ Add a random block to start
- ❑ Combine adjacent blocks to make ciphertext block
 - Many combination strategies



Cipher Block Chaining (CBC) mode encryption

AES

- ❑ Advanced Encryption Standard (2001)
 - Replaced DES (1977)
- ❑ Block size always 128 bits (4x4 bytes)
- ❑ Key size is 128, 192, or 256 bits
- ❑ Multi-step algorithm, many rounds



Security: Cryptography

Computer Science and Engineering ■ College of Engineering ■ The Ohio State University

Symmetric Key


- For ciphers (so far): Knowing E is enough to figure out D (its inverse)
 - If you know how to encrypt, you can decrypt too
 - Known as a *symmetric key* cipher
- Example: Caesar cipher
 - If $E(m) = m + 3$, $D(m) = m - 3$
- Example: One-time pad
 - Use same pad and same operation (xor)
- Example: AES
 - Use same key, reverse rounds and steps

One-Way Functions

- For some functions, the inverse is hard to calculate
 - One direction ($P \rightarrow Q$) is easy, but opposite direction ($Q \rightarrow P$) is hard/expensive/slow
- Intuition:
 - Given a puzzle *solution*, easy to design a puzzle with that solution (the “forward” direction)
 - Given the puzzle, hard to come up with the solution (the “inverse” direction)


Example: Dominating Set

- Hard direction: Find a dominating set of size at most 6 in the following graph...



Example: Dominating Set

- Easy direction: Create a graph with a dominating set of size 6 from this forest...



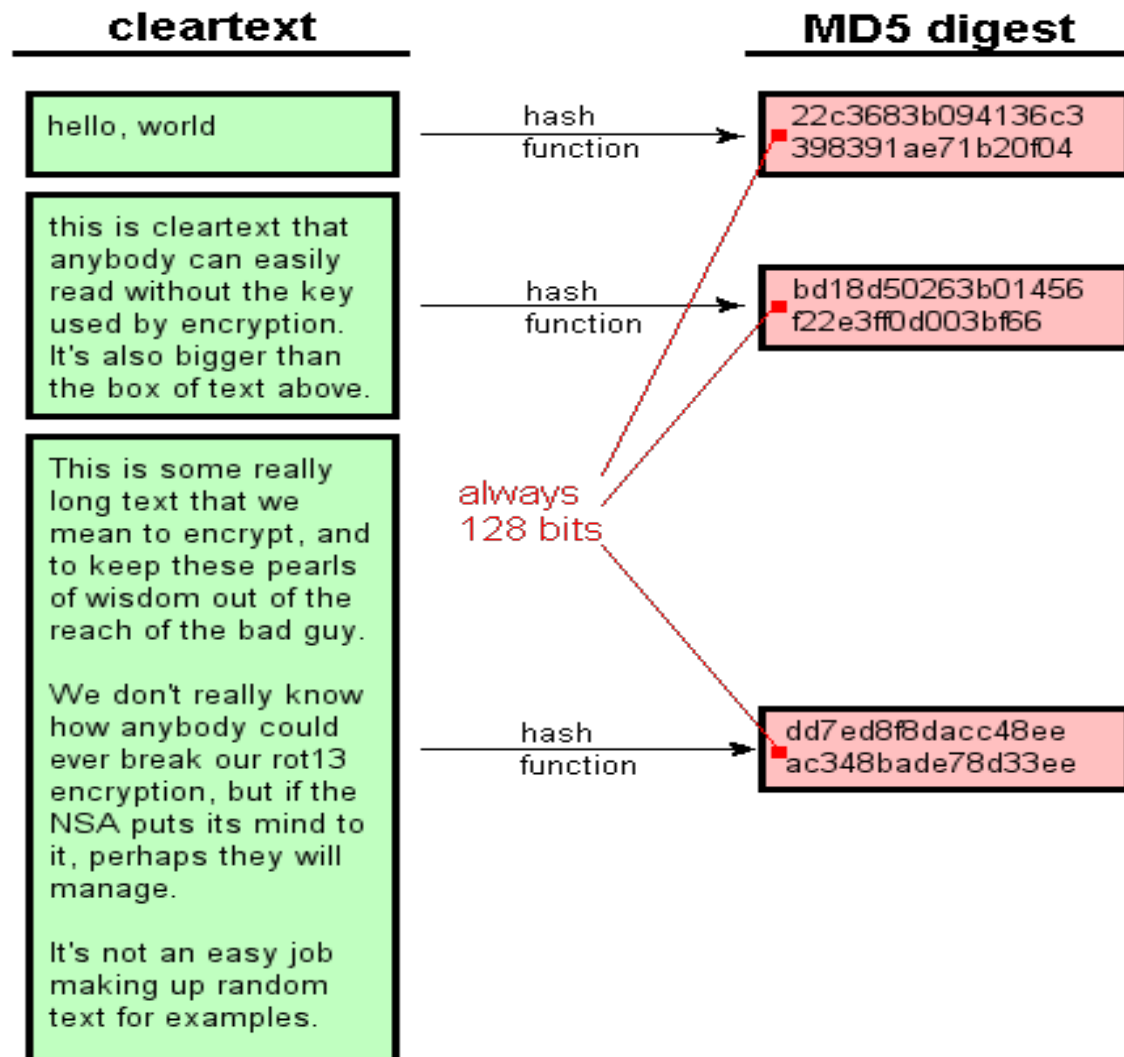
Example: Factoring

- Multiplying numbers is easy (*i.e.* fast)
 - Can multiply 2 n -bit numbers in n^2 steps
- Factoring a number is hard (*i.e.* slow)
 - To factor an n -bit number, need 2^n steps (approximately the number's value)
- Aside:
 - Primality testing is fast (recall lab activity in Software I and Fermat's Little Theorem)
 - But this fast test doesn't reveal the *factors* of a composite number

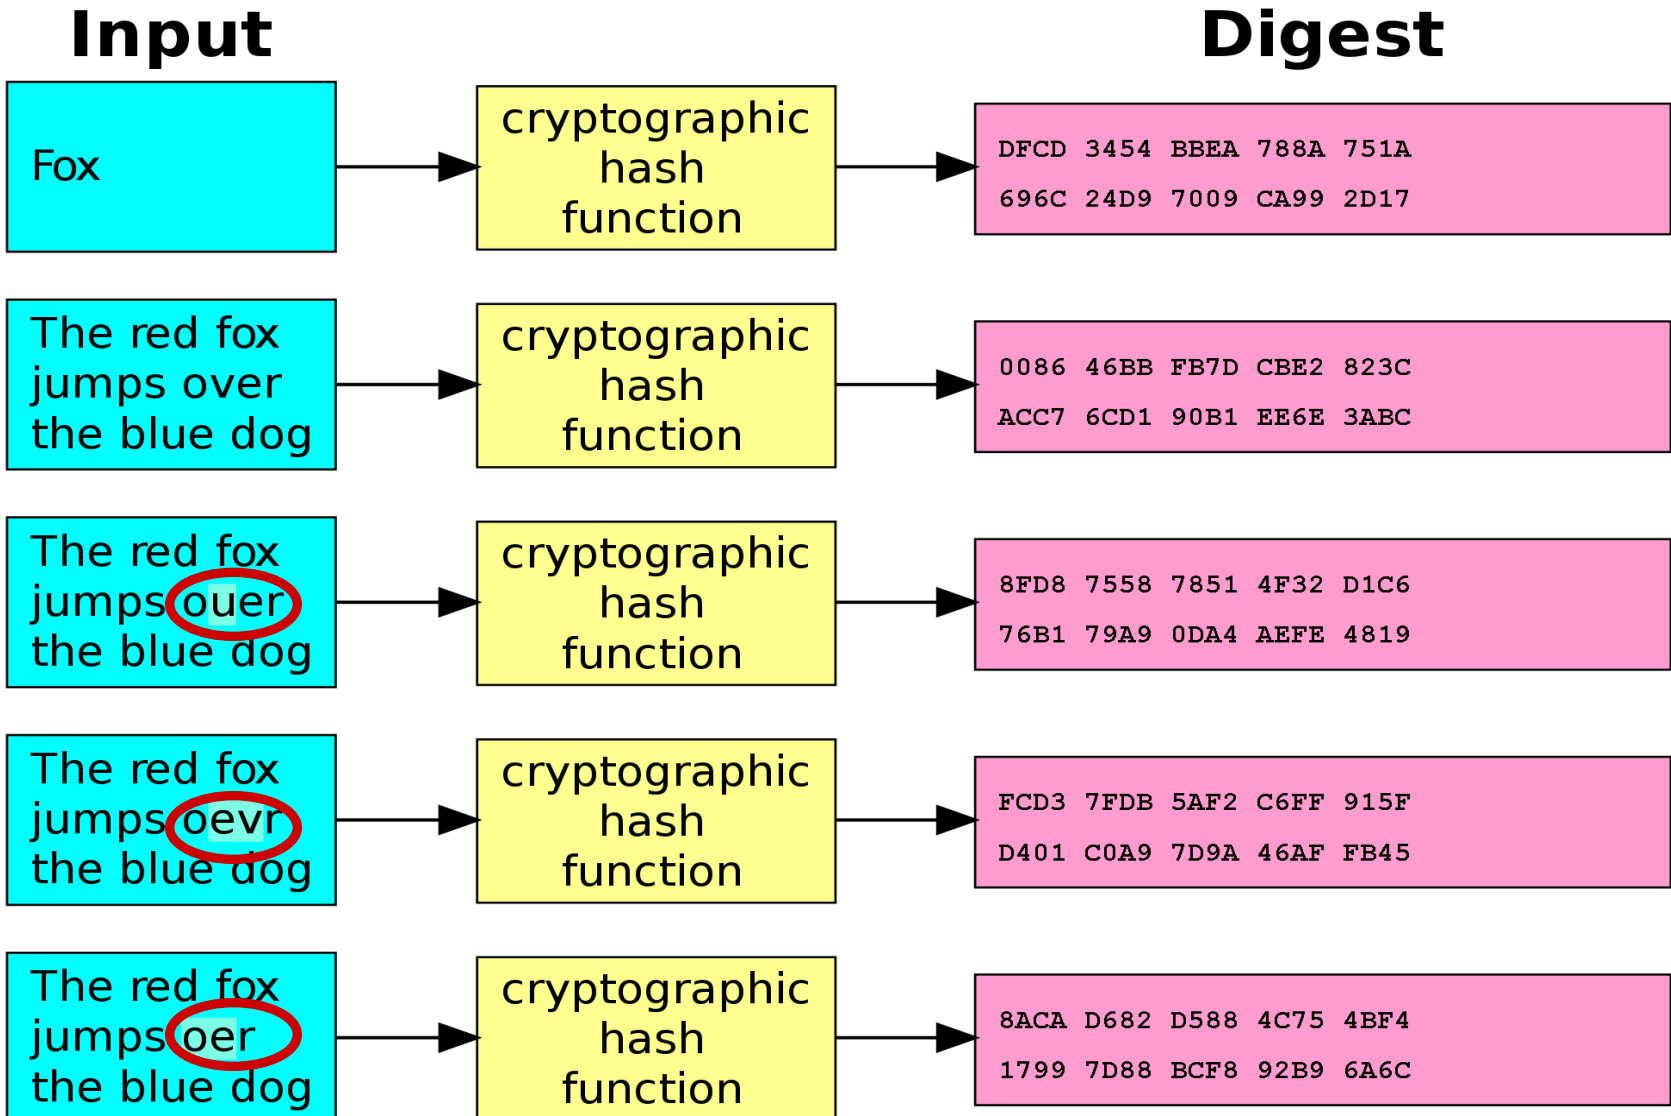
Cryptographic Hash Functions

- A hash function: $\mathbb{Z} \rightarrow \mathbb{Z}_B$
 - Every message, regardless of its length, maps to a number in the range $0..B - 1$
 - Result called a *digest* (constant-length, $\lg B$)
 - Good hashes give uniform distribution:
small diff in message \rightarrow big diff in digest
- *Cryptographic* hash func's are one-way
 - Given a digest, computationally infeasible to find *any* m that hashes to it
 - Collisions must still exist ($B \ll |\text{messages}|$), but are infeasible to find for large enough B
 - Digest = a *fingerprint* of m (small, fixed-size)

Fixed-Length Digests



Crypto. Hash as Fingerprint



Common Cryptographic Hashes

- MD5
 - Flaws discovered: “cryptographically broken”
 - Do not use!
- SHA-1: deprecated
 - Windows, Chrome, Firefox reject (2017)
 - 160-bit digests (*i.e.* 40 hex digits)
- Replaced by SHA-2 (still common)
 - A family of 6 different hash functions
 - Digest sizes: 224, 256, 384, or 512 bits
 - Names: **SHA**-224, SHA-256, SHA-512, etc
- Currently SHA-3
 - Entirely different algorithm
 - Names: **SHA3**-224, SHA3-256, SHA3-512, etc

Utility of Crypto. Hashes

- Integrity verification (super-checksum)
 - File download, check digest matches
- Password protection
 - Server stores the *hash* of user's password
 - Check entered password by computing its hash and comparing hash to the stored value
 - Benefit: Passwords are not stored (directly) in the database! If server is compromised, intruder finds hashes but not passwords
- Problem:
 - See md5decrypt.net/en/Sha256/
c023d5796452ad1d80263a05d11dc2a42b8c19c5d7c88c0e84ae3731b73a3d34

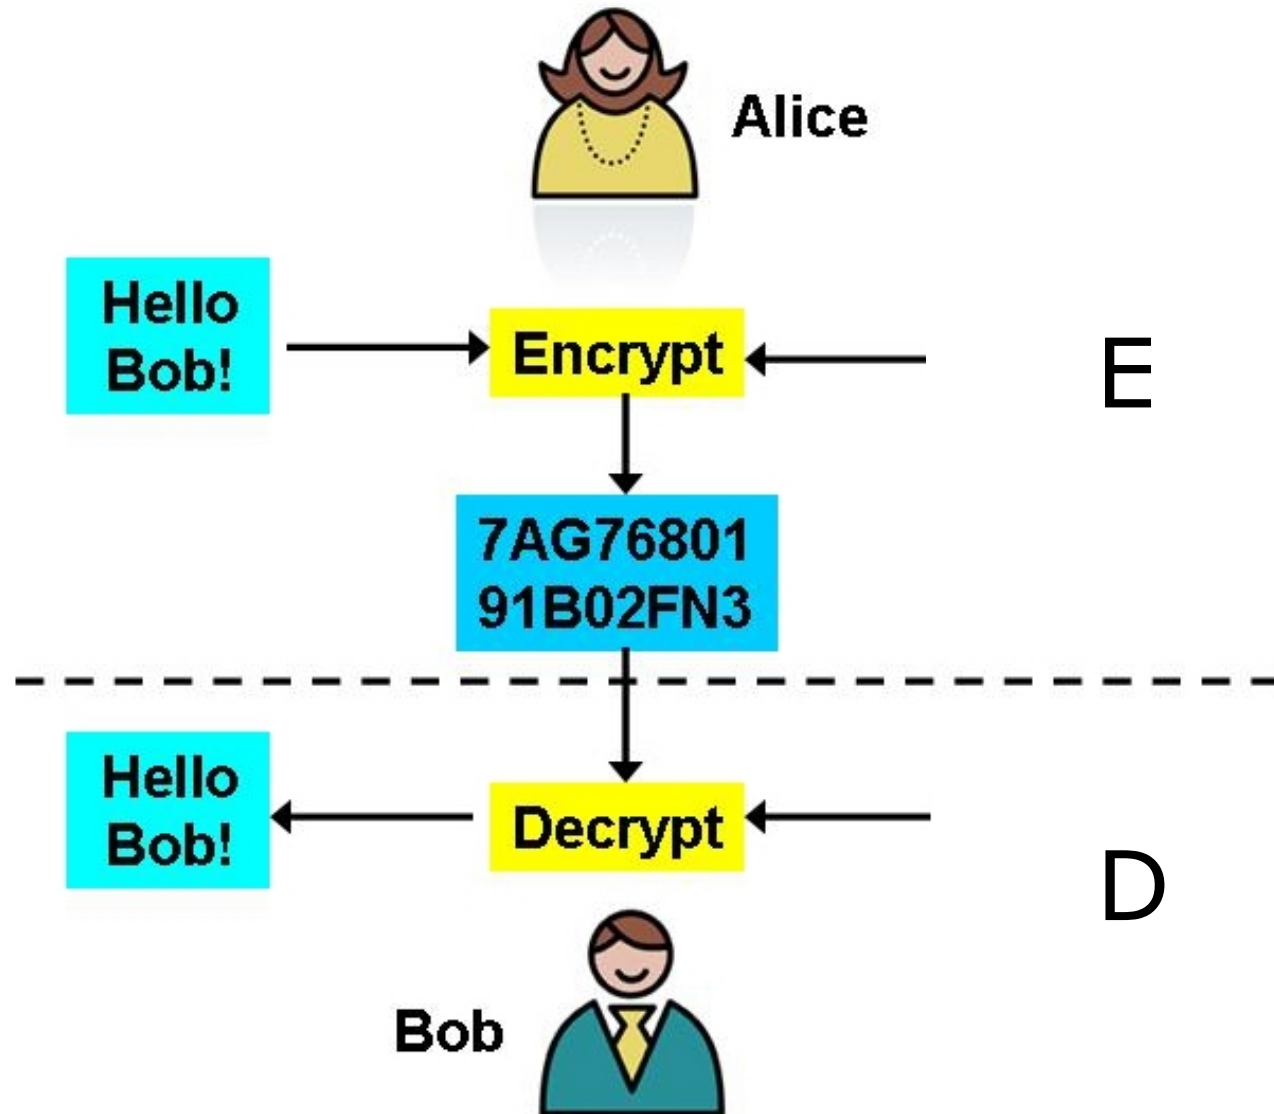
Role of Salt

- Danger:
 - Intruder pre-computes hashes for many (common) passwords: aka a *rainbow table*
 - Scan stolen hashes for matches
- Solution: *salt*
 - Server prepends text to password before hashing
 - Text must be *unique* to user
 - Text does not *need* to be secret
 - Ok: Deterministic value based on user name
 - Better: Random value, stored in the table
- Protects the fingerprint, by making it not mass pre-computable

One-Way Function with Trapdoor

- Function *appears* to be one-way
 - But, in reality, the inverse is easy if one knows a secret (the “trapdoor”)
- There are two very different functions:
 - The one-way-seeming function, E
 - The trapdoor for its inverse, D
- Knowing E is *not enough* to infer D
- Creates an asymmetry:
 - Alice knows E
 - Bob (and only Bob) knows D

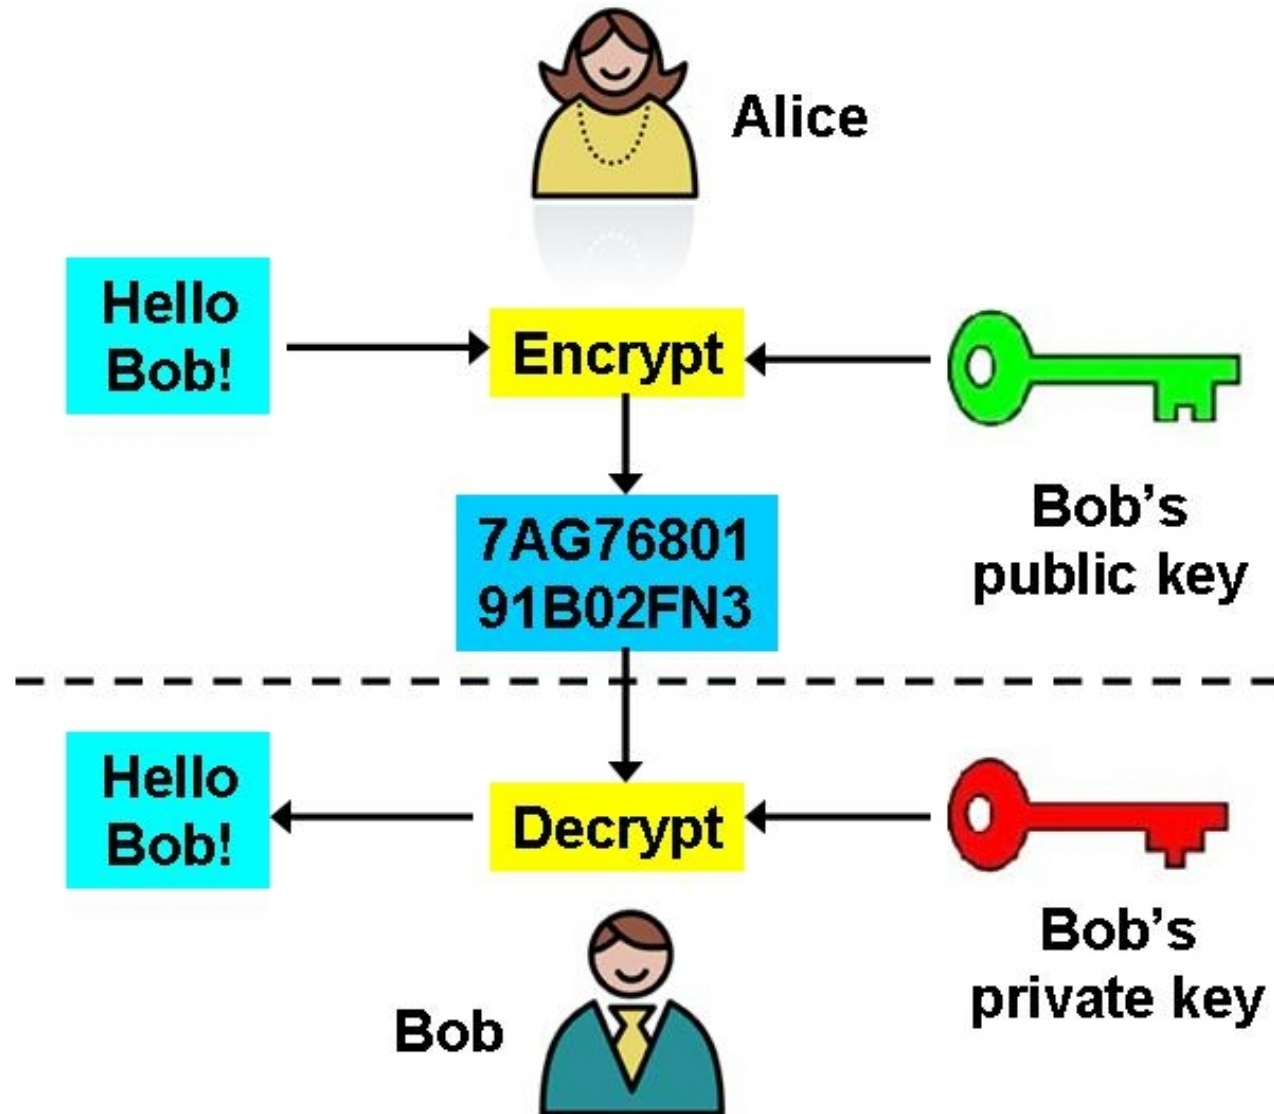
Asymmetry: Alice vs Bob



Public-Key Encryption

- Algorithms for E and D known by all
 - But parameterized by matched keys
- Asymmetry
 - Key for Bob's E is *public*
 - Key for Bob's D is *private*
- Anyone can encrypt messages for Bob
- Only Bob can decrypt these messages
- Important consequences
 - Each agent needs only 1 public key
 - No pre-existing shared secret needed

Public and Private Keys



RSA

- E and D are actually the same function

$$m^k \bmod n$$

- Parameterized by pair (k, n) , i.e. the key

- Private key: (d, n)

- $D(m) = m^d \bmod n$

- Public key: (e, n)

- $E(m) = m^e \bmod n$

- Choice of e & d is based on factoring

- Choose 2 large **prime** numbers, p and q

- Calculate their product, $n = pq$

- Pick any d relatively prime with $(p-1)(q-1)$

- Find an e s.t. $ed = 1 \bmod (p-1)(q-1)$

Digital Signature

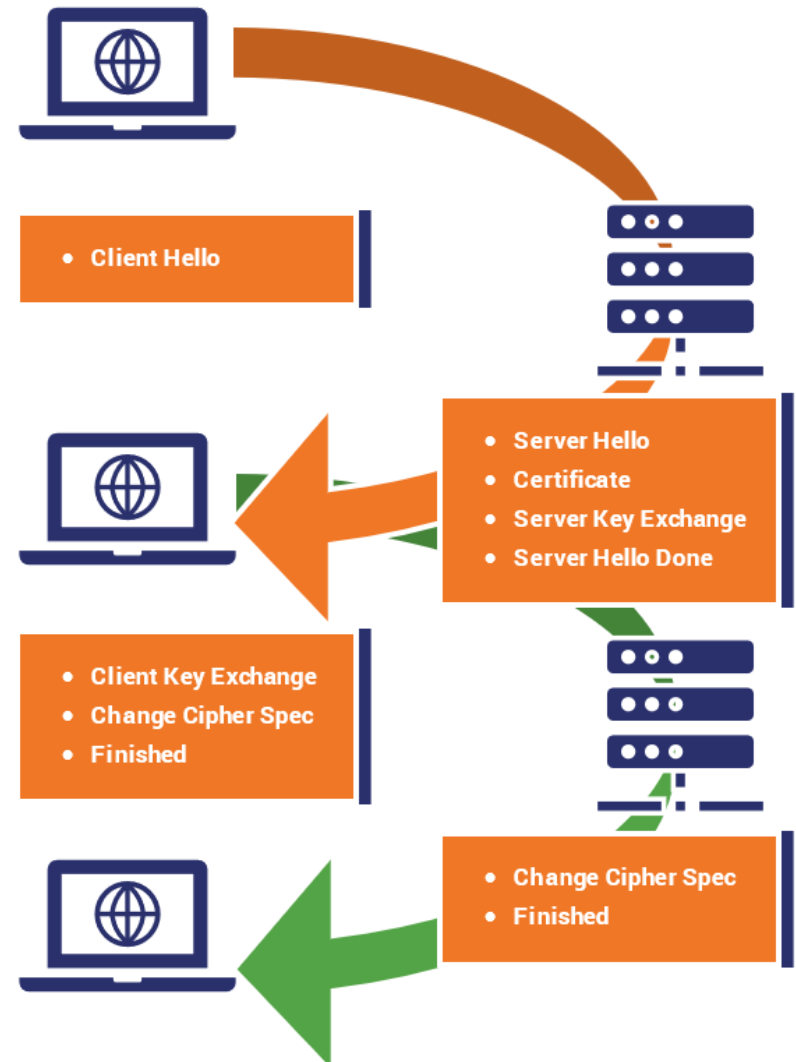
- Usual direction for encryption:
$$D(E(m)) = (m^e)^d = m^{ed} = m, \text{ mod } n$$
- One-to-one, so backwards works too!
$$E(D(m)) = (m^d)^e = m^{de} = m, \text{ mod } n$$
- Consider:
 - Bob “encrypts” m using his **private** key, d
 - Bob sends **both** m and $D(m)$
 - Anyone can undo the “encrypted” part using Bob’s **public** key, e
 - Result will be m
- $D(m)$ serves as a digital **signature** of m
 - Only Bob could have created this signature
 - Use: non-repudiation

Performance Considerations

- ❑ Symmetric key algorithms are faster than public key algorithms
- ❑ Optimization for encryption (RSA)
 - Create a fresh symmetric key, k
 - Use symmetric algorithm to encrypt m
 - Use recipient's public key to encrypt k
- ❑ Optimization for digital signatures
 - Calculate the digest for m (always short)
 - Use sender's private key to encrypt digest

TLS 1.3: Handshake

- Certificate authority
 - Connects public key to identity
- Client:
 - Get server's public key
 - Make new (symmetric) session key
 - Sends this key to server (encrypted with public key)



Take Home Message

- ❑ Don't try to roll your own crypto/security implementation
- ❑ Use (trusted) libraries
- ❑ Recognize role and importance of (eg):
 - Initialization vector
 - Cryptographic hash/digest
 - Salt
 - Private key vs public key

Summary

- Symmetric-key encryption
 - Sender and receiver share (same) secret key
 - Stream ciphers work one bit at a time (*e.g.*, one-time pad)
 - Block ciphers work on larger blocks of bits (*e.g.*, SHA-2)
- One-way functions: Hard to invert
 - Cryptographic hash produces fixed-size digest
 - Digest serves as a fingerprint
- Public key encryption
 - Matching keys: k_{private} , k_{public}
 - Anyone can use public key to encrypt
 - Only holder of private key can decrypt
 - Use private key to create a digital signature